

Recent Progress in and around LibreSSL

Theo Buehler
tb@openbsd.org

EuroBSDCon – September 17, 2022

About LibreSSL

One of the four major forks of OpenSSL

- ▶ 1998: OpenSSL forks from/continues SSLeay
accumulates (more) shoddy code, cruft over next 16 years
after lots of disasters, heartbleed makes people look and act
- ▶ Apr 2014: OpenBSD forks LibreSSL
- ▶ Jun 2014: Adam Langley (Google) makes BoringSSL public
- ▶ Nov 2021: Akamai / Microsoft want QUIC \rightsquigarrow QuicTLS
OpenSSL + patchset to add BoringSSL QUIC API

LibreSSL Main Features

- ▶ libtls: sane and easy-to-use wrapper of the SSL/TLS stack
- ▶ clean room implementation of TLSv1.3 stack (2018-2020)
 - ▶ centerpieces: record layer and handshake state machine
 - ▶ missing features: PSK (work in progress), ECH (complicated)
 - ▶ non-goal: early data
- ▶ new certificate validator
- ▶ documentation (unfortunately there's only one schwarze@)
- ▶ lots of code cleanup
- ▶ largely compatible with OpenSSL 1.1 on support intersection
- ▶ this improved a lot due to making structs in LibreSSL opaque
- ▶ ABI about as stable as OpenSSL 1.1

On OpenSSL compatibility

- ▶ OpenSSL 1.1 API: have what we need, more than we wanted
- ▶ No OpenSSL 3 API yet
- ▶ > 2000 OpenBSD ports link against libcrypto or libssl
- ▶ < 100 of these need patches (< 5%)
- ▶ Painful: Qt, PyPy (because of py-cryptography), stunnel
- ▶ By far the most requested missing feature is Ed25519 ...
- ▶ ... followed by things like SHA-512/256, SHA-3, Blake, ...
- ▶ 6 ports link against OpenSSL:
 - ▶ mail/opensmtpd-filters/dkimsign flavor (Ed25519 signatures)
 - ▶ mail/postfix (DANE, mostly)
 - ▶ net/bro aka zeek: needs TLS-PRF API
 - ▶ lang/node: Ed25519 + a dozen API functions
 - ▶ net/nagios/nsca-ng: PSK
 - ▶ security/libretls: by design

Background: Anatomy of a Certificate

Certificates are a complicated data structure.

ASN.1 from RFC 5280:

```
Certificate ::= SEQUENCE {  
    tbsCertificate      TBSCertificate,  
    signatureAlgorithm  AlgorithmIdentifier,  
    signatureValue      BIT STRING }
```

- ▶ sequence: basically a struct
- ▶ TBS: To Be Signed
- ▶ Contents of struct
 1. what is (to be) signed
 2. how is it signed
 3. signature

Background: Anatomy of a Certificate (continued)

```
TBSCertificate ::= SEQUENCE {  
    version          [0] EXPLICIT Version DEFAULT v1,  
    serialNumber      CertificateSerialNumber,  
    signature         AlgorithmIdentifier,  
    issuer            Name,  
    validity          Validity,  
    subject           Name,  
    subjectPublicKeyInfo SubjectPublicKeyInfo,  
    issuerUniqueID    [1] IMPLICIT UniqueIdentifier OPTIONAL,  
                      -- If present, version MUST be v2 or v3  
    subjectUniqueID   [2] IMPLICIT UniqueIdentifier OPTIONAL,  
                      -- If present, version MUST be v2 or v3  
    extensions        [3] EXPLICIT Extensions OPTIONAL  
                      -- If present, version MUST be v3  
}
```

Could go on forever.

RFC 5280: 151 pages

> 80 of which are the details of this struct (and CRLs)

A PEM encoded certificate

Most of you will have seen something like this

```
-----BEGIN CERTIFICATE-----
```

```
MIIG4jCCBcqqgAwIBAgISBJxsswkXm1b9UVavEeONm1EzMA0GCSqGSIb3DQFB  
MDIxCzAJBgNVBAYTA1VTMRYwFAYDVQQKEw1MZXQncyBFbmNyeXBOMQswCQYD  
EwJSMzAeFw0yMjA3MjExNTQ3MTJaFw0yMjEwMTkxNTQ3MTFaMB0xGDAWBgl
```

```
...
```

```
-----END CERTIFICATE-----
```

- ▶ PEM: Privacy Enhanced Mail (see RFC 7468)
- ▶ Base64 encoded DER of certificate
- ▶ DER: Distinguished Encoding (Rules) of ASN.1 “struct”

Aside: Why do certs start with MII?

All 133 CA certs in OpenBSD's root bundle start with MII

```
$ grep -c -- -----BEGIN /etc/ssl/cert.pem
```

```
133
```

```
$ grep -A1 -- -----BEGIN /etc/ssl/cert.pem | grep -c MII
```

```
133
```

```
$ echo -n MIIG | b64decode -r | hexdump -Cv | head -n 1  
00000000  30 82 06
```

- ▶ 30: DER: encoding of an ASN.1 SEQUENCE
- ▶ 82: DER: the length is described by the next two bytes
- ▶ MII: Base64 of 30 82 + 2 most significant bits of length
- ▶ length of a cert is > 127 bytes (so needs at least two bytes)
- ▶ length of a cert is usually < 16684 bytes,
so the two most significant bits are 0

New Certificate Validator

- ▶ “Legacy validator” inherited from OpenSSL: unmaintainable
- ▶ During lockdown, beck@ wrote an RFC 5280 validator
- ▶ Initial code was correct. We only found minor bugs, ...
- ▶ ... then many months of whack-a-mole started
- ▶ Lots of software relies on
 - ▶ strange and overly specific error codes in certain situations
 - ▶ undocumented behavior of the verify callback
 - ▶ specific order of traversing the potential chains
- ▶ Took us two years to be reasonably compatible with the legacy validator
 - ▶ fix one thing, break ten others
 - ▶ one hole introduced in the process

Legacy Record Layer Rewrite (WIP)

- ▶ jsing@ wrote a very nice record layer underlying TLSv1.3
- ▶ Similar ideas can be used for old TLS versions and DTLS
- ▶ Goal: remove `ssl_pkt.c` and `d1_pkt.c` (terrible code)
- ▶ Uses CBS and CBB instead of explicit pointer manipulations
- ▶ With this work, DTLSv1.2 support came pretty much for free
 - ▶ landry@: linphone, baresip
 - ▶ kn@: tdesktop
 - ▶ missing bit: BIO_ADDR API, so Qt cannot yet use it

QUIC API

- ▶ De facto standard API by David Benjamin of BoringSSL
- ▶ OpenSSL PR 8797 (2019): port by Todd Short (Akamai)
- ▶ Had to wait for OpenSSL 3 (was already late at that point)
- ▶ Mai 2021: QUIC standardized in RFCs 9000 – 9002
- ▶ Sep 2021: OpenSSL 3 released
- ▶ Oct 2021: OpenSSL want their own stack
 - ▶ BoringSSL compatibility explicit non-goal
 - ▶ Unclear why. *Someone* must have a reason. . .
 - ▶ QUIC transport protocol not really within OpenSSL's expertise
- ▶ Nov 2021: QuicTLS announced in IETF side meeting

QUIC API (continued)

- ▶ beck@ and jsing@ ported BoringSSL API
- ▶ Plugged very nicely into jsing@'s record layer
- ▶ Needed `EVP_chacha20_poly1305` support in libcrypto
- ▶ Experimental version will be available in LibreSSL 3.6
 - ▶ curl can speak QUIC using ngtcp2
 - ▶ wlallemand added minimal working version to haproxy
 - Needs `SSL_CTX_set_client_hello_cb` for full support
- ▶ BoringSSL API works, but is not great
 - ▶ exposes full structs and enums publically (sigh...)
 - ▶ BoringSSL and QuicTLS have already diverged
 - ▶ ngtcp2 initializes public struct without C99 initializers
 - ▶ BoringSSL open to improvements
 - ▶ QuicTLS probably set in stone

Primality Testing

Starting point: a 2018 preprint:

Prime and Prejudice: Primality Testing Under Adversarial Conditions.

Albrecht, Massimo, Paterson, Somorovsky:

- ▶ [...] construct 2048-bit composites that are declared prime with probability $1/16$
- ▶ [...] the advertised performance [LibreSSL/OpenSSL] is 2^{-80}
- ▶ [...] for a number of libraries (Cryptlib, LibTomCrypt, JavaScript Big Number, WolfSSL), we can construct composites that always pass the supplied primality tests

Primality Testing (continued)

Tricky to fix

- ▶ Workaround: crank number of Miller-Rabin rounds (slow)
- ▶ Recommendation: Baillie–Pomerance–Selfridge–Wagstaff algorithm
- ▶ Problem: this isn't easy – someone needs time and skills

Primality Testing (continued)

Lucky coincidence: Martin Grenouilloux has time and skills

- ▶ background: espie@ finds preprint independently
- ▶ tells us he has a promising student with a knack for maths
- ▶ Martin already had a Python implementation
- ▶ a few weeks later: C implementation lands in my inbox
- ▶ work stalled for a few of weeks due to exams
- ▶ things become easier with a mostly correct implementation...
- ▶ clean up, optimize, simplify, fix, and commit
- ▶ result is one of the nicest pieces of code in libcrypto
- ▶ amazing work by Martin Grenouilloux

RFC 3779 support

- ▶ This is about routing and BGP
- ▶ X.509 Extensions for IP Addresses and AS Identifiers
- ▶ Issuer of certificate transfers “internet numbers” to subject
- ▶ Part of libcrypto, ported by job@ from OpenSSL
- ▶ Helps `rpki-client`, makes `openssl x509` output nicer
- ▶ Needed audit, cleanup, lots of fixes, regress
- ▶ Public API is pretty broken
- ▶ Downside: code is inefficient, hit by certificate validator
- ▶ `rpki-client`: spends $\sim 10\%$ of runtime in RFC 3779 code

Testing, CI and Coverity

- ▶ Ilya Shipitsin from haproxy has been tremendously helpful
 - ▶ Helped add ASAN CI, which has been invaluable
 - ▶ Also helps with triaging Coverity issues
- ▶ tlsfuzzer runs as part of daily regression tests
 - ▶ Tickles many corner cases
 - ▶ Helped improve standards compliance a lot
 - ▶ Hannes Mehnert mentioned it at BSDCan 2019, thanks!
- ▶ The Ruby OpenSSL Gem has a very useful test suite
- ▶ Joshua Sing rewrote and improved many of the old tests

Thanks

- ▶ LibreSSL core team: bcook@, beck@, inoguchi@, jsing@
- ▶ schwarze@ for awesome documentation and many bug fixes
- ▶ ajacoutot@, sthen@ for help with ports
- ▶ genua@ for testing infrastructure and for sponsoring work
- ▶ Martin Grenouilloux, espie@ for the work on primality testing
- ▶ Ilya Shipitsin for help with portable
- ▶ “orbea” for helping with upstream patches
- ▶ OpenBSD foundation for sponsoring bulk build machine